



Crime Audio Event Classification

Final project for *Digital Forensics and Biometrics*

Luca Domeneghetti (mat. 2159079)

University of Padova A.A. 2024/2025



Project description

The project required a Neural Network model to be built and trained in order to perform a **classification** task on audio samples, mostly concerning *crime events*.

In total, the model shall perform classification between 13 audio classes:

car_crash	conversation	engine_idling	gun_shot
jambret	maling	rain	rampok
road_traffic	scream	thunderstorm	tolong
wind			



Datasets





- **Raw** audio of accident and crime detection
 - Around 1600 audio samples, each with variable duration
 - <https://www.kaggle.com/datasets/afisarsy/raw-audio-of-accident-and-crime-detection>
- **Enhanced** audio of accident and crime detection
 - More than 9000 audio samples
 - Generated by combining the raw dataset samples with noise effects (wind, thunderstorm, rain and road traffic)
 - Allows for a better generalization of audio features
 - <https://www.kaggle.com/datasets/afisarsy/enhanced-audio-of-accident-and-crime-detection>

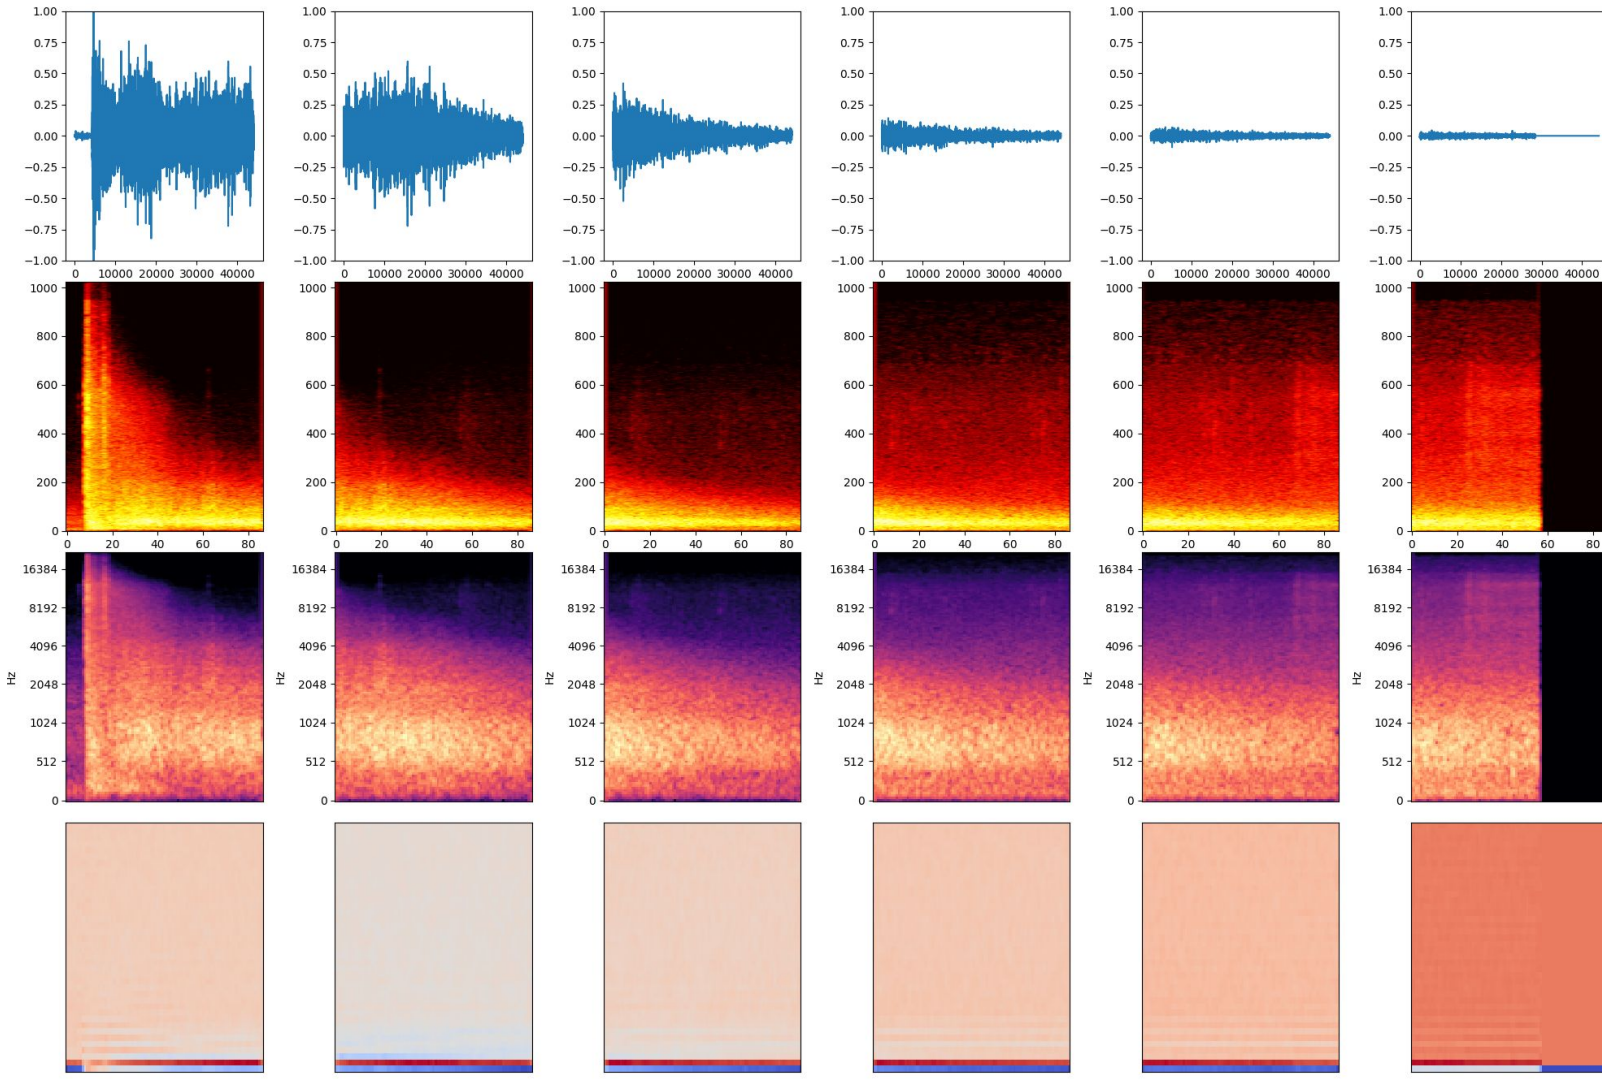


Audio data preprocessing

- **Chunking** was adopted to split the longer audio files into smaller frames
 - *Partial overlap* to avoid information loss on the frame's boundaries.
- **Mel spectrogram**
 - Operate on the frequency domain instead of the time domain
- **Mel Frequency cepstral coefficients** further enhanced the capability of the model to extract hidden features from data
 - Both the Mel spectrogram and MFCC computations were performed within the LSTM network in order to better exploit the performance of the GPU







Neural Network and Framework

- Initial approaches involved the use of **Convolutional Neural Networks**:
 - **1D CNN** were used on a flattened vector made out of both the Mel spectrogram and the MFCC combined.
 - **2D CNN** were used to classify spectrograms and MFCC in a “image-recognition” fashion.
 - Although they are easier to understand, CNN failed to properly grasp a fundamental aspect of audio data: time correlation and dependencies.
- For such reasons, it was chosen to abandon CNNs to use some other architectures.


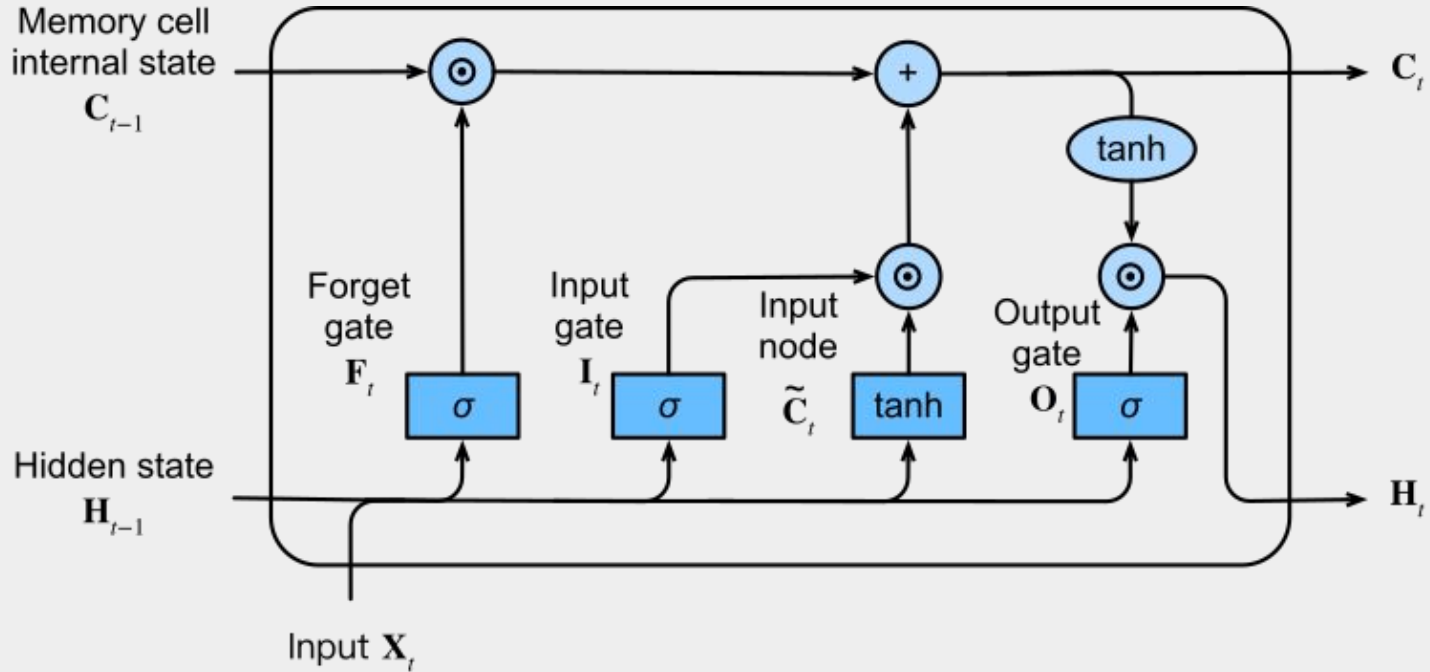


Neural Network and Framework

- The core of the project is a **LSTM RNN**
 - RNN is a good choice when it comes to time-based data.
 - LSTM network is efficient at solving some issues related with the vanishing/exploding of the gradient.
- Initially TensorFlow
 - TensorFlow simpler at handling the dataset and construction of LSTM network, but lacked the required CUDA support.
- Switched to **PyTorch**
 - PyTorch needed additional learning time, but fitted the scope of running the training process entirely on the GPU.




LSTM neuron




FC layer with
activation function



Elementwise
operator



Copy



Concatenate

LSTM Network

- Internal **Mel/MFCC** feature computation:
 - 128 Mel bands
 - 40 MFCC coefficients
- Two **bidirectional LSTM** with a single hidden layer of 256 units
- Combine **last two hidden layers** of the forward and backward layers of the LSTM
- Three **dense ReLU layers** with a **dropout** (0.5) in between the first and the last two
- Use of **batch normalization**
- Return **logits** value



Training function

- Loss: **cross entropy** applied on *sparse categorical* predictions against *one-hot encoded labels*
- Optimizer: **AdamW** with *learning rate* = 0.001 and *weight decay* = $1e-4$
- Scheduler: **OneCycleLR** to better fit learning rate to learning conditions



First training attempt

- Use of **raw dataset**
- **Frame size** of 0.5 seconds
- **Hop size** of 0.2 seconds
- **Batches** of 16 audio samples
- **16 epochs**
- Training on Nvidia GeForce RTX 3060 with CUDA 12.7
- Training duration: ~10 minutes

Training accuracy	Test accuracy
87%	71%



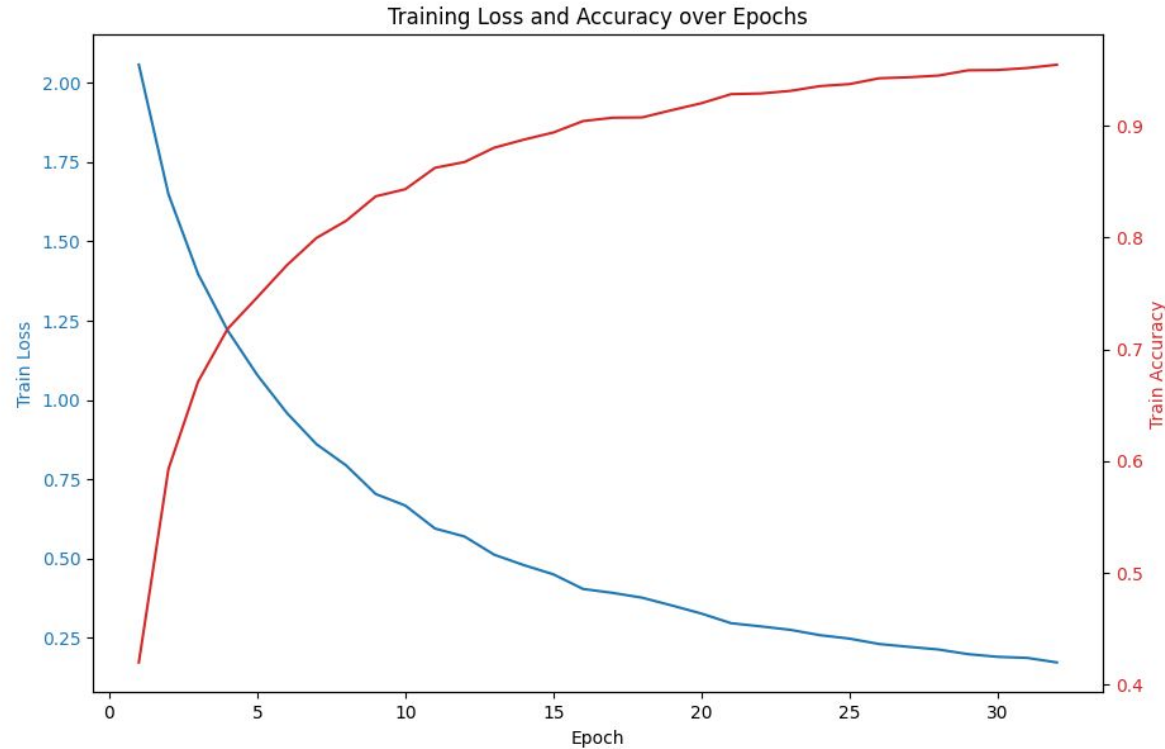
Final training results

- Use of **enhanced dataset**
- **Frame size** of 0.4 seconds
- **Hop size** of 0.2 seconds
- **Batches** of 32 audio samples
- **32 epochs**
- Training on Nvidia GeForce RTX 3060 with CUDA 12.7
- Training duration: ~15 minutes


Training accuracy	Test accuracy
95.45%	94.92%



Training graph



Confusion Matrix



Prediction

- Similar fashion to audio pre-processing:
 - Split full sample into smaller windows with partial overlap
 - Perform prediction on each window
 - Collate predictions to form “event segments”
 - Compute start/finish instant in seconds and print
- Samples taken from online sound effect databases
 - samplefocus.com pixabay.com myinstant.com
 - Complex audio samples created by merging together simpler samples using Audacity, resampling to 44.1 kHz, converting to .wav



Bibliography

- [seth814/Audio-Classification](#), for the structure of the LSTM network
- [Audio Classification Starter](#), for the use of Datasets and DataLoaders
- [Audio Classification with LSTM](#), for the training and testing function

